

# LOAN AMOUT PREDICTION

In [1]:

```
# import Libraries
```

```
import numpy as np # Linear algebra
import pandas as pd # data processing
import os
for dirname, _, filenames in os.walk('/jaya/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

In [2]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
%matplotlib inline
```

In [3]:

```
df = pd.read_csv("train_data.csv")
df
```

Out[3]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	Loan_Status
0	LP001002	Male	No	0	Graduate	No	5849		Approved
1	LP001003	Male	Yes	1	Graduate	No	4583		Approved
2	LP001005	Male	Yes	0	Graduate	Yes	3000		Approved
3	LP001006	Male	Yes	0	Not Graduate	No	2583		Rejected
4	LP001008	Male	No	0	Graduate	No	6000		Approved
...	...	...	...	...	...	...	...	...	...
609	LP002978	Female	No	0	Graduate	No	2900		Approved
610	LP002979	Male	Yes	3+	Graduate	No	4106		Approved
611	LP002983	Male	Yes	1	Graduate	No	8072		Approved
612	LP002984	Male	Yes	2	Graduate	No	7583		Approved
613	LP002990	Female	No	0	Graduate	Yes	4583		Approved

614 rows × 13 columns

In [4]:

```
df.head()
# 1st 5 rows of data
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

In [5]: df.tail()# Last 5 rows of data

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
609	LP002978	Female	No	0	Graduate	No	2900	
610	LP002979	Male	Yes	3+	Graduate	No	4106	
611	LP002983	Male	Yes	1	Graduate	No	8072	
612	LP002984	Male	Yes	2	Graduate	No	7583	
613	LP002990	Female	No	0	Graduate	Yes	4583	

In [6]: print("No.of Rows and Columns:",df.shape)

No.of Rows and Columns: (614, 13)

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   Loan_ID          614 non-null    object  
 1   Gender           601 non-null    object  
 2   Married          611 non-null    object  
 3   Dependents       599 non-null    object  
 4   Education        614 non-null    object  
 5   Self_Employed    582 non-null    object  
 6   ApplicantIncome  614 non-null    int64  
 7   CoapplicantIncome 614 non-null    float64 
 8   LoanAmount        592 non-null    float64 
 9   Loan_Amount_Term  600 non-null    float64 
 10  Credit_History   564 non-null    float64 
 11  Property_Area    614 non-null    object  
 12  Loan_Status       614 non-null    object  
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

In [8]: df.describe()

Out[8]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
<b>count</b>	614.000000	614.000000	592.000000	600.000000	564.000000
<b>mean</b>	5403.459283	1621.245798	146.412162	342.000000	0.842199
<b>std</b>	6109.041673	2926.248369	85.587325	65.12041	0.364878
<b>min</b>	150.000000	0.000000	9.000000	12.00000	0.000000
<b>25%</b>	2877.500000	0.000000	100.000000	360.000000	1.000000
<b>50%</b>	3812.500000	1188.500000	128.000000	360.000000	1.000000
<b>75%</b>	5795.000000	2297.250000	168.000000	360.000000	1.000000
<b>max</b>	81000.000000	41667.000000	700.000000	480.000000	1.000000

In [9]: df.isnull().sum()

```
Out[9]:
```

Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0

dtype: int64

## Analysis of each columns and also Data cleaning

Data cleaning removing null values replacing duplicate values

In [10]: df['Gender'].value\_counts(dropna=False)

```
Out[10]:
```

Male	489
Female	112
NaN	13

Name: Gender, dtype: int64

```
In [11]: # Here we can replace the null values with Male
df['Gender']=df['Gender'].replace(np.nan,'Male')
```

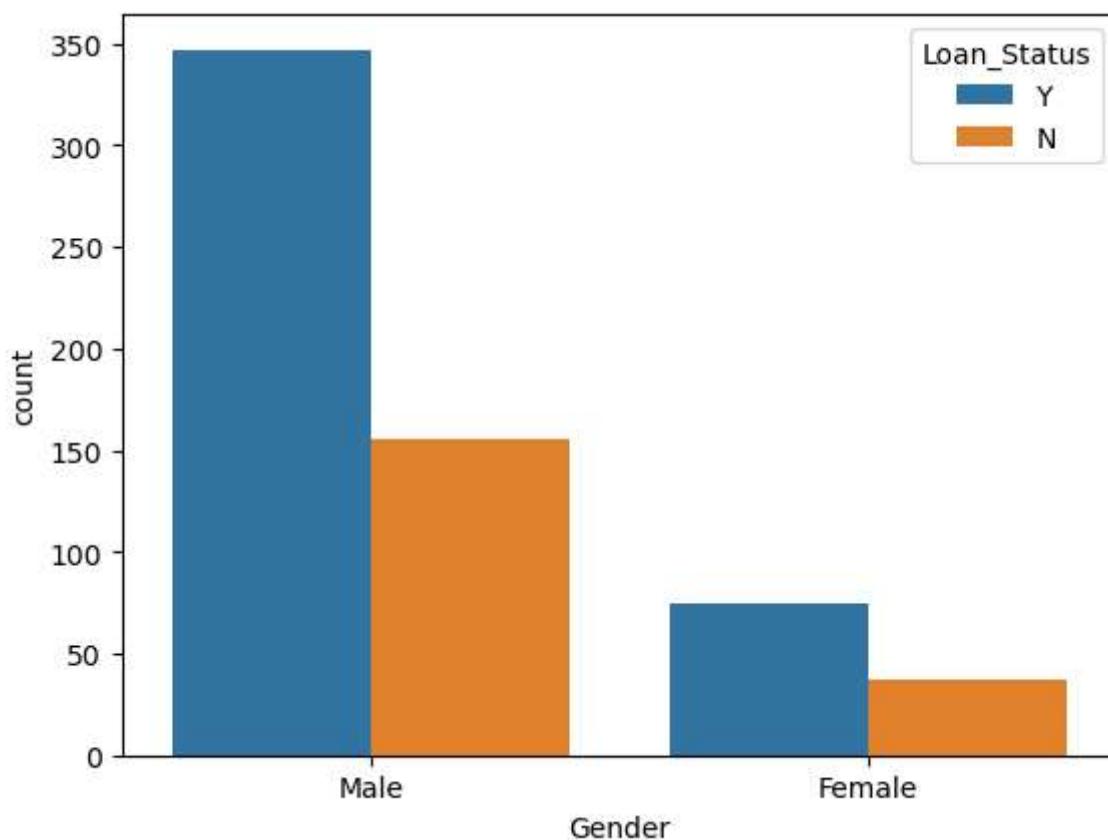
## visulization

using countplot

In [12]: # COUNT PLOTE USING SEABORN IN GENDER COLUMN

```
sns.countplot(x=df['Gender'],hue=df['Loan_Status'])
```

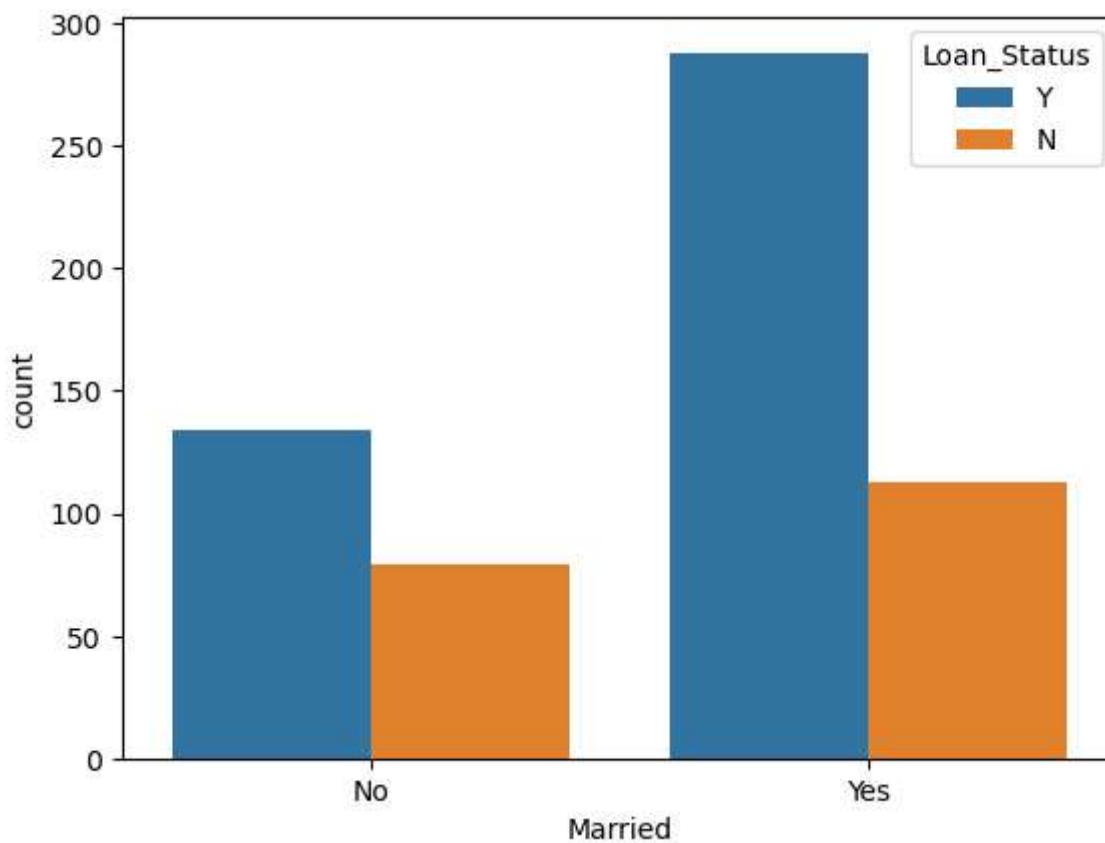
Out[12]: &lt;AxesSubplot:xlabel='Gender', ylabel='count'&gt;

In [13]: `df['Married'].value_counts(dropna=False)`

Out[13]:

In [14]: `# here we replace the null values with Yes  
df['Married']=df['Married'].replace(np.nan,'Yes')`In [15]: `sns.countplot(x=df['Married'],hue=df['Loan_Status'])`

Out[15]: &lt;AxesSubplot:xlabel='Married', ylabel='count'&gt;



we can observe married applicants have more chances of getting loan

```
In [16]: df['Dependents'].value_counts(dropna=False)
```

```
Out[16]:
0    345
1    102
2    101
3+   51
NaN  15
Name: Dependents, dtype: int64
```

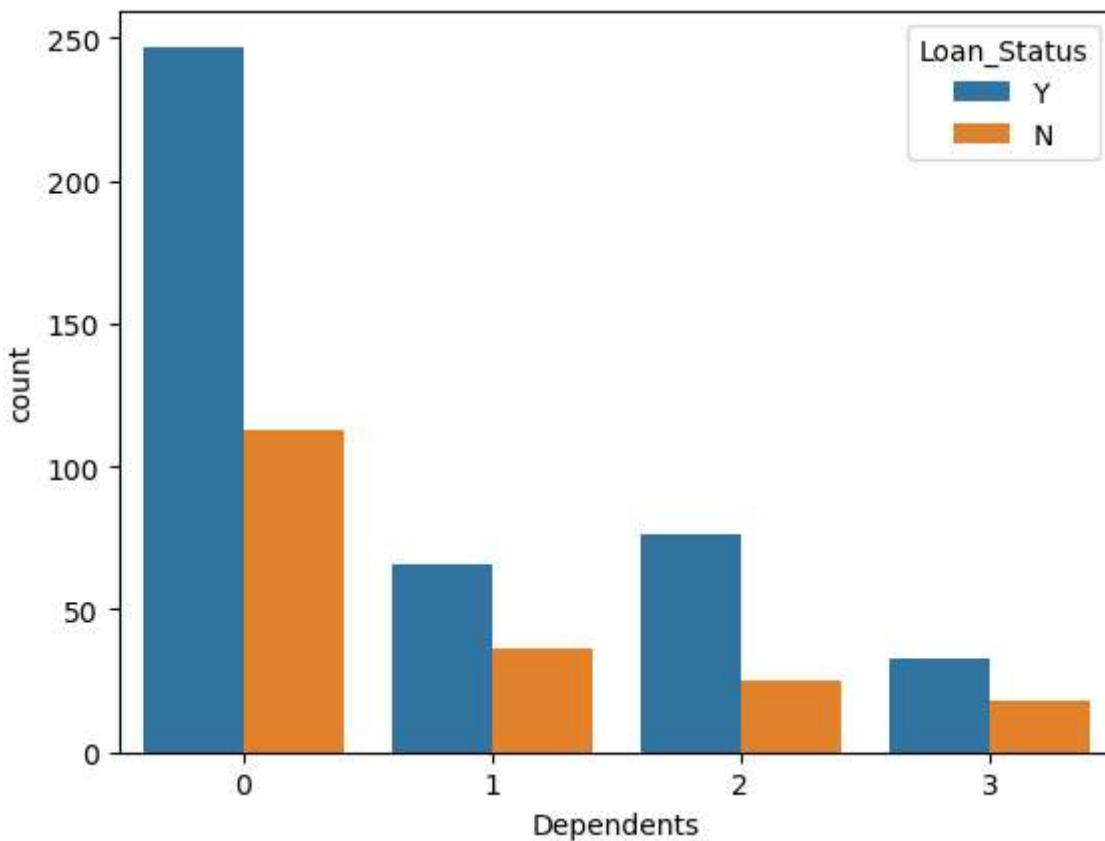
```
In [17]: df['Dependents']=df['Dependents'].replace(np.nan,'0')
df['Dependents']=df['Dependents'].replace('3+', '3')
```

```
In [18]: df['Dependents'].value_counts(dropna=False)
```

```
Out[18]:
0    360
1    102
2    101
3    51
Name: Dependents, dtype: int64
```

```
In [19]: sns.countplot(x=df['Dependents'],hue=df['Loan_Status'])
```

```
Out[19]: <AxesSubplot:xlabel='Dependents', ylabel='count'>
```

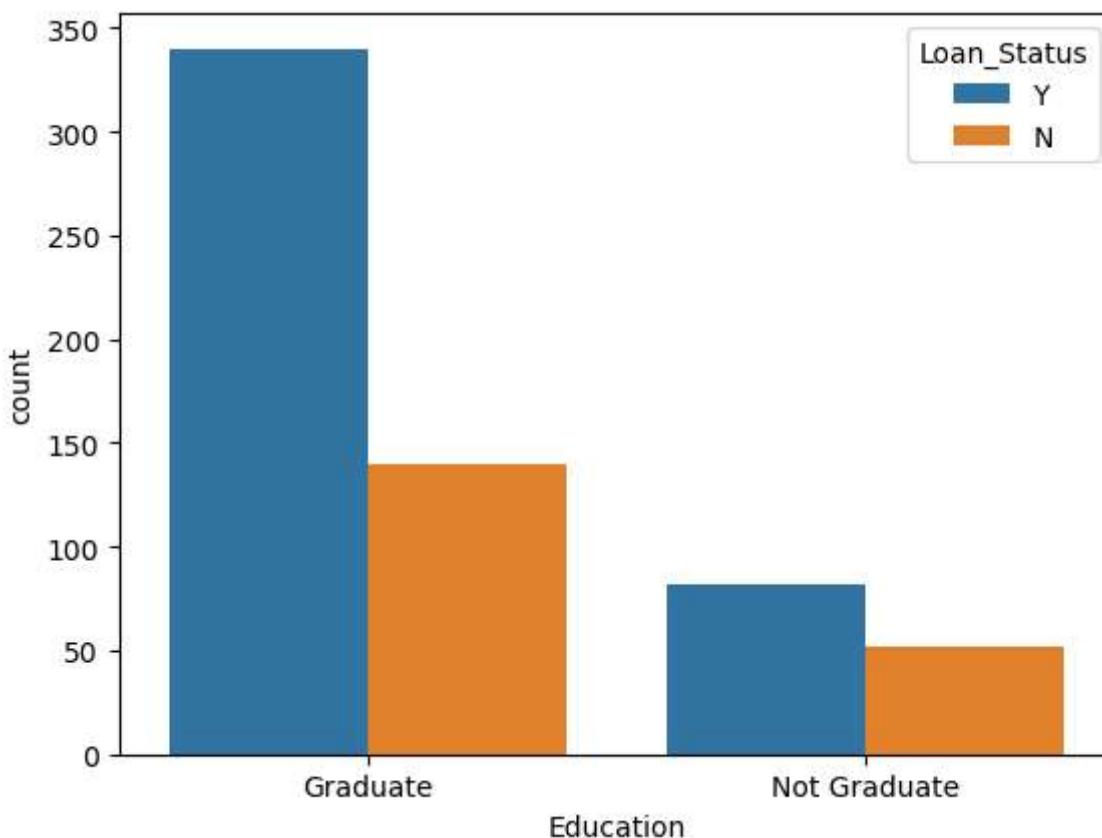


```
In [20]: df['Education'].value_counts(dropna=False)
```

```
Out[20]: Graduate      480  
Not Graduate   134  
Name: Education, dtype: int64
```

```
In [21]: sns.countplot(x=df['Education'], hue=df['Loan_Status'])
```

```
Out[21]: <AxesSubplot:xlabel='Education', ylabel='count'>
```



we can see graduates have higher chances of getting loan

```
In [22]: df['Self_Employed'].value_counts(dropna=False)
```

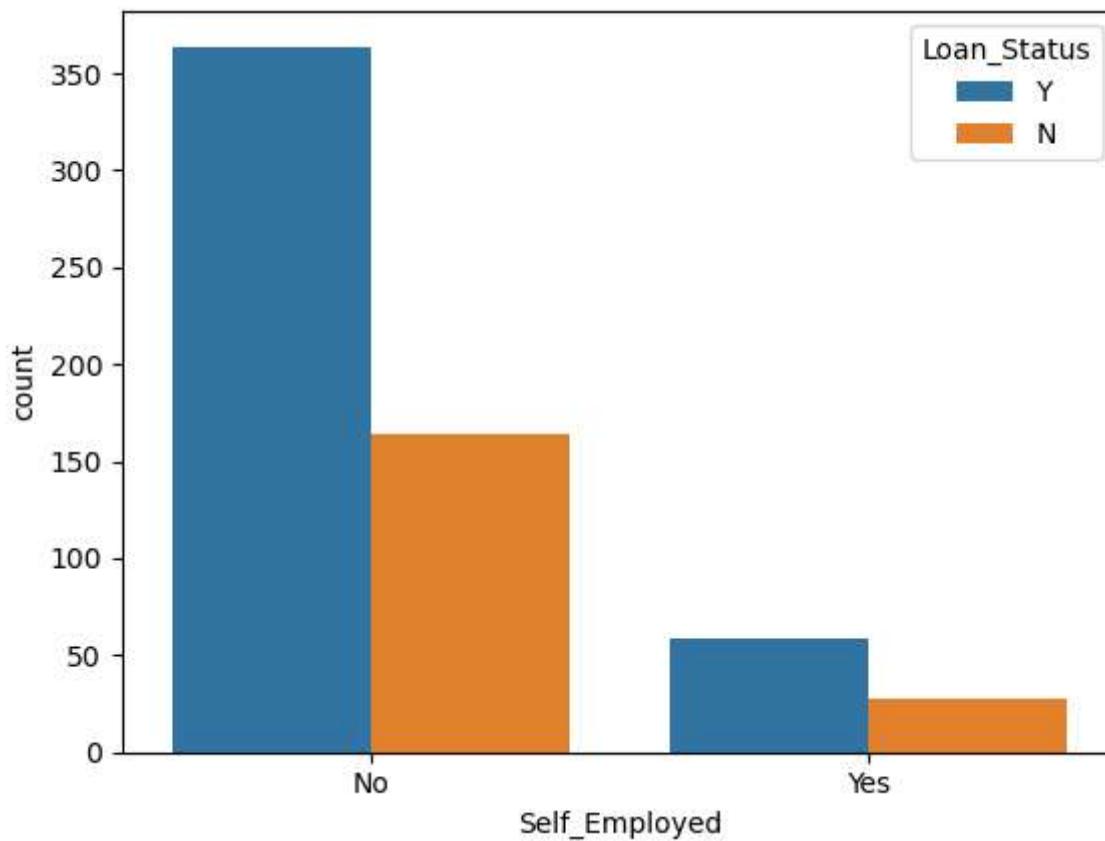
```
Out[22]: No      500
          Yes     82
          NaN     32
          Name: Self_Employed, dtype: int64
```

## bfill used for filling null values

```
In [23]: # bfill used for filling null values
df['Self_Employed']=df['Self_Employed'].bfill()
```

```
In [24]: sns.countplot(x=df['Self_Employed'],hue=df['Loan_Status'])
```

```
Out[24]: <AxesSubplot:xlabel='Self_Employed', ylabel='count'>
```



for self employed customers have less chance of loan approval

```
In [25]: df['ApplicantIncome'].value_counts(dropna=False)
```

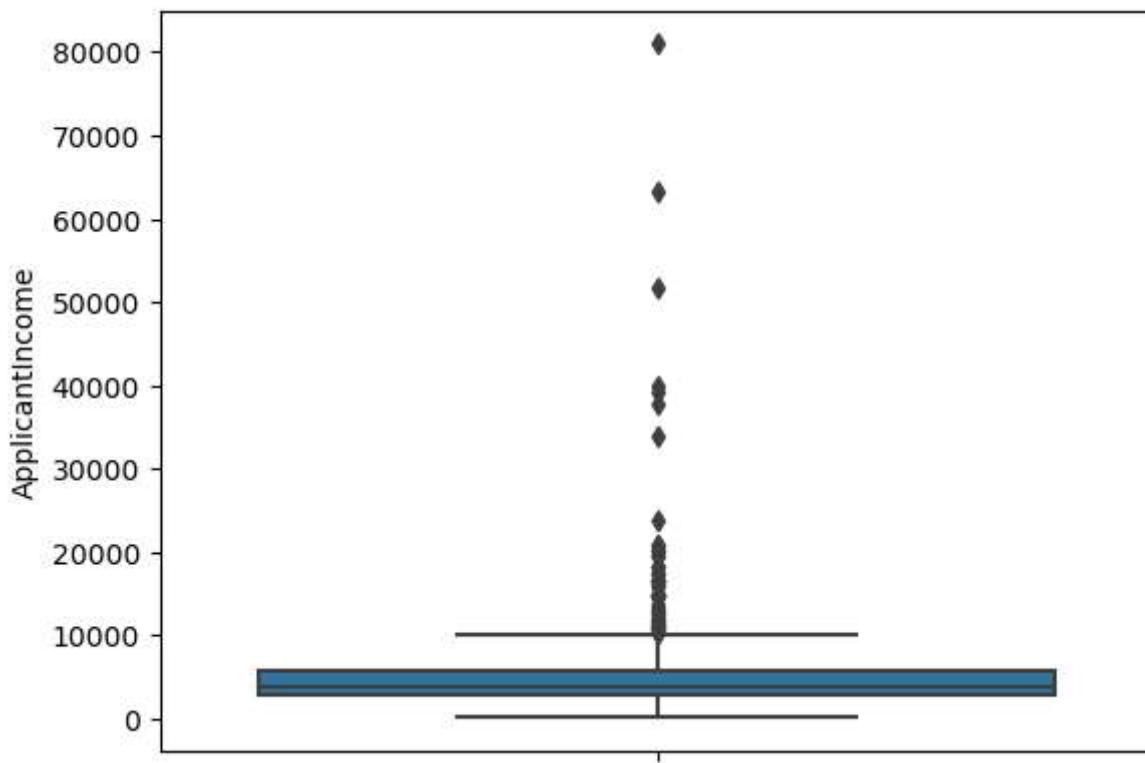
```
Out[25]:
```

2500	9
4583	6
6000	6
2600	6
3333	5
..	
3244	1
4408	1
3917	1
3992	1
7583	1

Name: ApplicantIncome, Length: 505, dtype: int64

```
In [26]: #visualizing ApplicantIncome numeric variable  
sns.boxplot(y=df['ApplicantIncome'])
```

```
Out[26]: <AxesSubplot:ylabel='ApplicantIncome'>
```



```
In [27]: #checking percentile values
df['ApplicantIncome'].describe(percentiles=(0.01,0.05,0.1,0.25,0.5,0.75,0.9,0.95,0.99))
```

```
Out[27]: count      614.000000
mean      5403.459283
std       6109.041673
min      150.000000
1%       1025.000000
5%       1897.550000
10%      2216.100000
25%      2877.500000
50%      3812.500000
75%      5795.000000
90%      9459.900000
95%     14583.000000
99%     32540.410000
max      81000.000000
Name: ApplicantIncome, dtype: float64
```

## capping

capping is an outlier technic

uses for removing outliers

removes top&bottom 1 %

```
In [28]: Q1=df['ApplicantIncome'].quantile(0.01)
Q3=df['ApplicantIncome'].quantile(0.99)
```

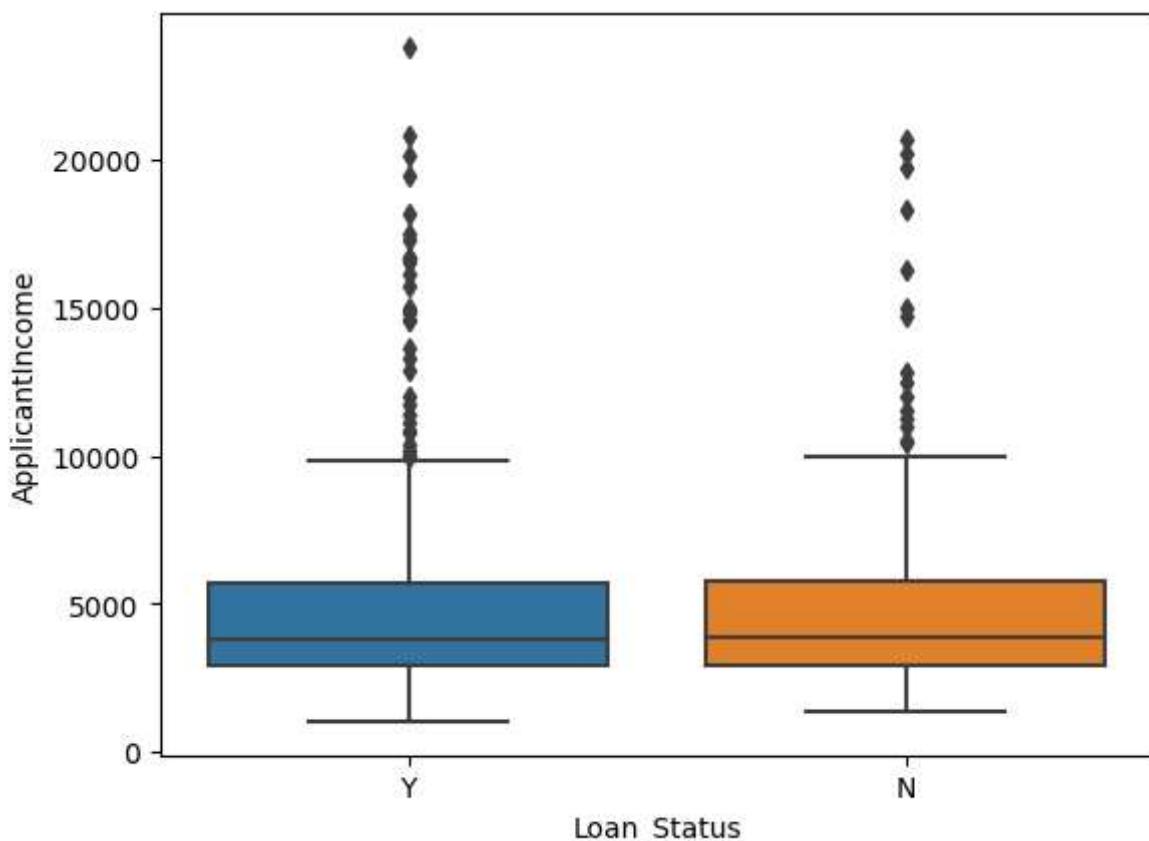
```
In [29]: df=df[df['ApplicantIncome']<=Q3]
df=df[df['ApplicantIncome']>=Q1]
```

```
In [30]: df.shape
```

```
Out[30]: (601, 13)
```

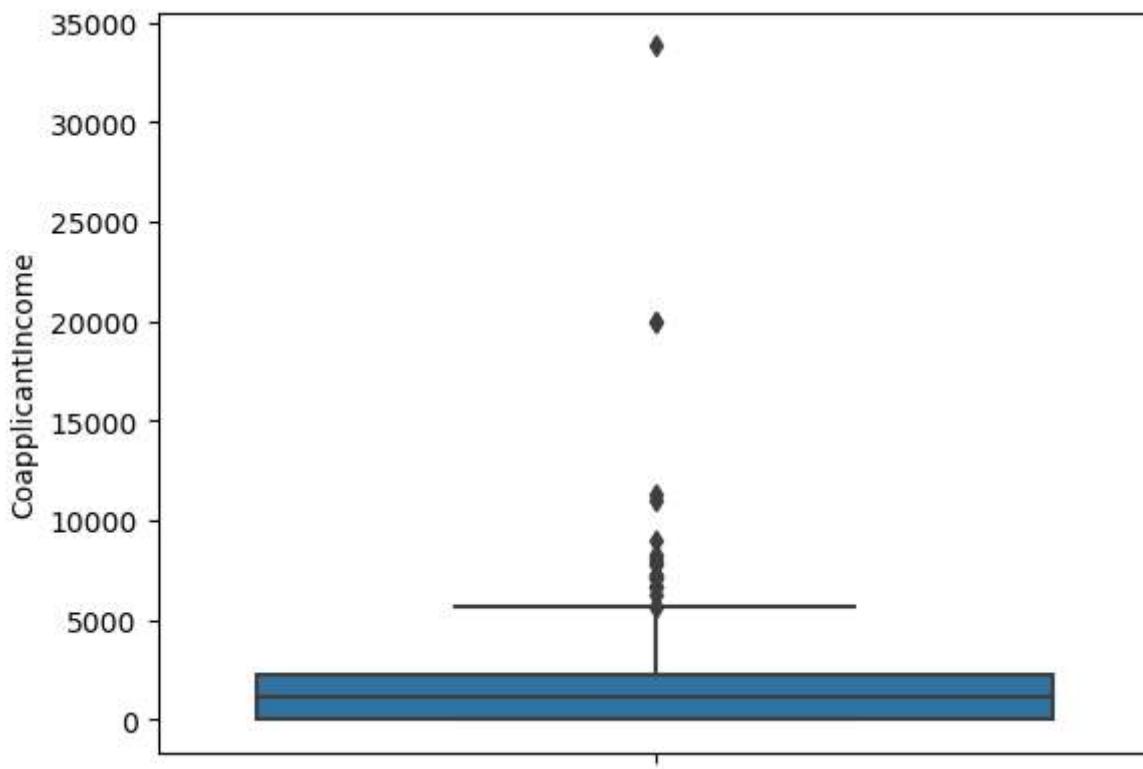
```
In [31]: #checking Spread of ApplicantIncome vs Loan_Status variable
sns.boxplot(y=df['ApplicantIncome'],x=df['Loan_Status'])
```

```
Out[31]: <AxesSubplot:xlabel='Loan_Status', ylabel='ApplicantIncome'>
```



```
In [32]: #visualizing spread of CoapplicantIncome numeric variable
sns.boxplot(y=df['CoapplicantIncome'])
```

```
Out[32]: <AxesSubplot:ylabel='CoapplicantIncome'>
```



```
In [33]: #checking percentiles for CoapplicantIncome
df['CoapplicantIncome'].describe(percentiles=(0.01,0.05,0.1,0.25,0.5,0.75,0.9,0.95,0.99))
```

```
Out[33]:
count      601.000000
mean      1551.264426
std       2447.952413
min       0.000000
1%        0.000000
5%        0.000000
10%       0.000000
25%       0.000000
50%       1167.000000
75%       2254.000000
90%       3750.000000
95%       4983.000000
99%       8333.000000
max      33837.000000
Name: CoapplicantIncome, dtype: float64
```

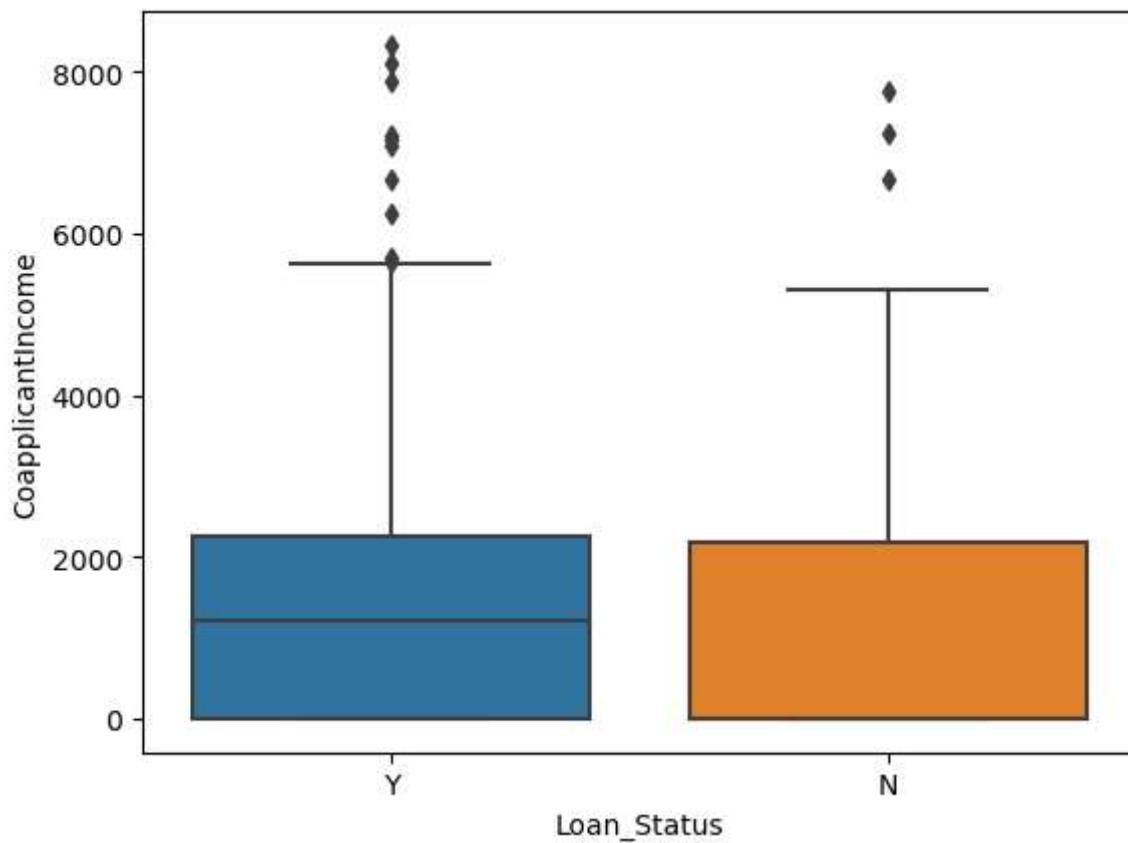
```
In [34]: #outlier treatment - Remove top&bottom 1%
Q1=df['CoapplicantIncome'].quantile(0.01)
Q3=df['CoapplicantIncome'].quantile(0.99)
df=df[df['CoapplicantIncome']<=Q3]
df=df[df['CoapplicantIncome']>=Q1]
```

```
In [35]: df.shape
```

```
Out[35]: (595, 13)
```

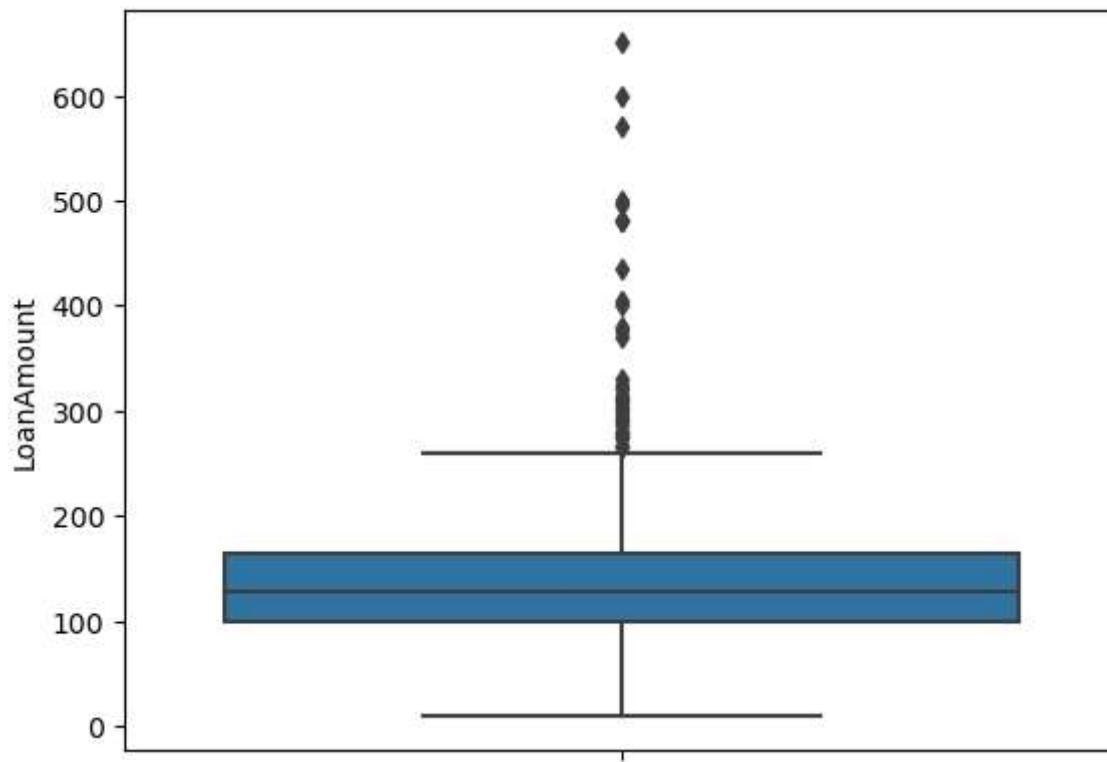
```
In [36]: #boxplot of CoapplicantIncome vs Loan_Status variable
sns.boxplot(y=df['CoapplicantIncome'],x=df['Loan_Status'])
```

```
Out[36]: <AxesSubplot:xlabel='Loan_Status', ylabel='CoapplicantIncome'>
```



```
In [37]: #visualizing spread of LoanAmount numeric variable  
sns.boxplot(y=df['LoanAmount'])
```

```
Out[37]: <AxesSubplot:ylabel='LoanAmount'>
```



```
In [38]: df['LoanAmount'].isnull().sum()
```

Out[38]: 22

```
In [39]: #checking percentiles for "LoanAmount"
df['LoanAmount'].describe(percentiles=(0.01,0.05,0.1,0.25,0.5,0.75,0.9,0.95,0.99))
```

```
Out[39]: count    573.000000
mean     142.420593
std      77.182008
min      9.000000
1%       30.000000
5%       55.600000
10%      71.000000
25%      100.000000
50%      127.000000
75%      165.000000
90%      217.600000
95%      276.600000
99%      480.000000
max      650.000000
Name: LoanAmount, dtype: float64
```

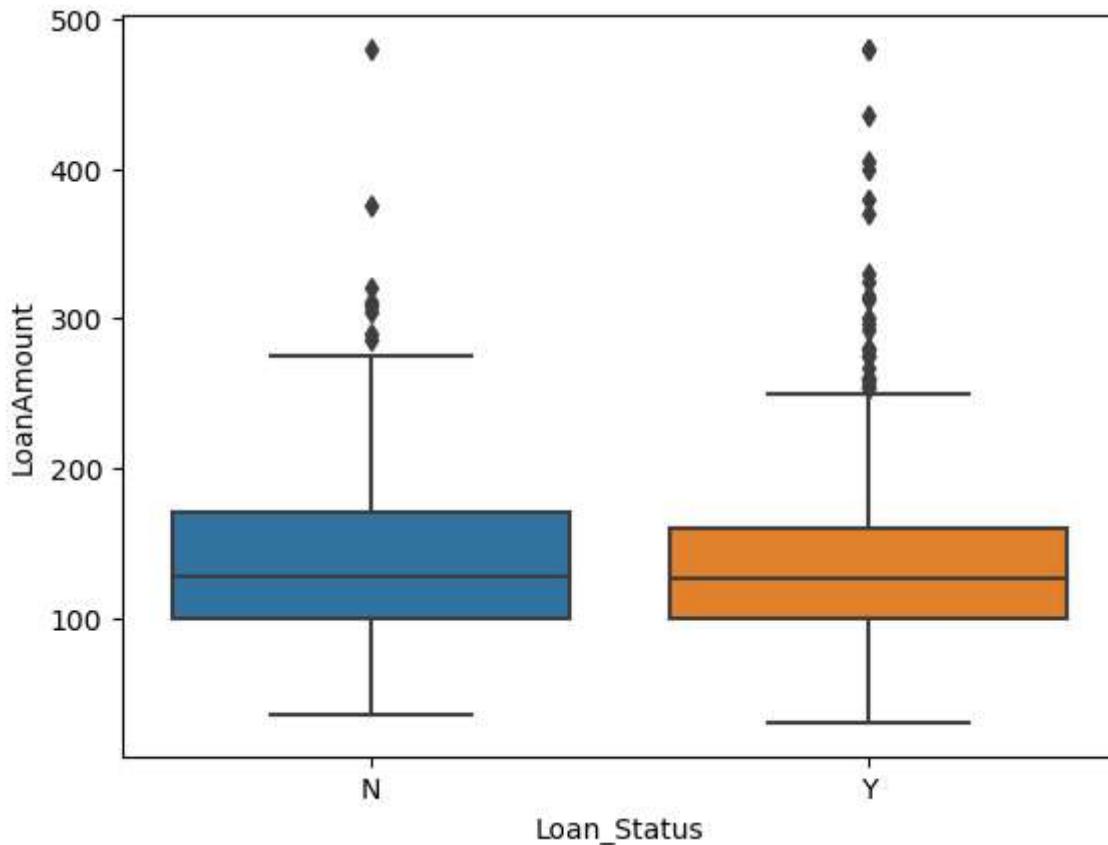
```
In [40]: #Outlier Treatment: Remove top & bottom 1%
Q1=df['LoanAmount'].quantile(0.01)
Q3=df['LoanAmount'].quantile(0.99)
df=df[df['LoanAmount']<=Q3]
df=df[df['LoanAmount']>=Q1]
```

In [41]: df.shape

Out[41]: (563, 13)

```
In [42]: #checking Spread of "LoanAmount" vs Loan_Status variable
sns.boxplot(y=df['LoanAmount'],x=df['Loan_Status'])
```

Out[42]: &lt;AxesSubplot:xlabel='Loan\_Status', ylabel='LoanAmount'&gt;



```
In [43]: df['Loan_Amount_Term'].value_counts(dropna=False)
```

```
Out[43]:
```

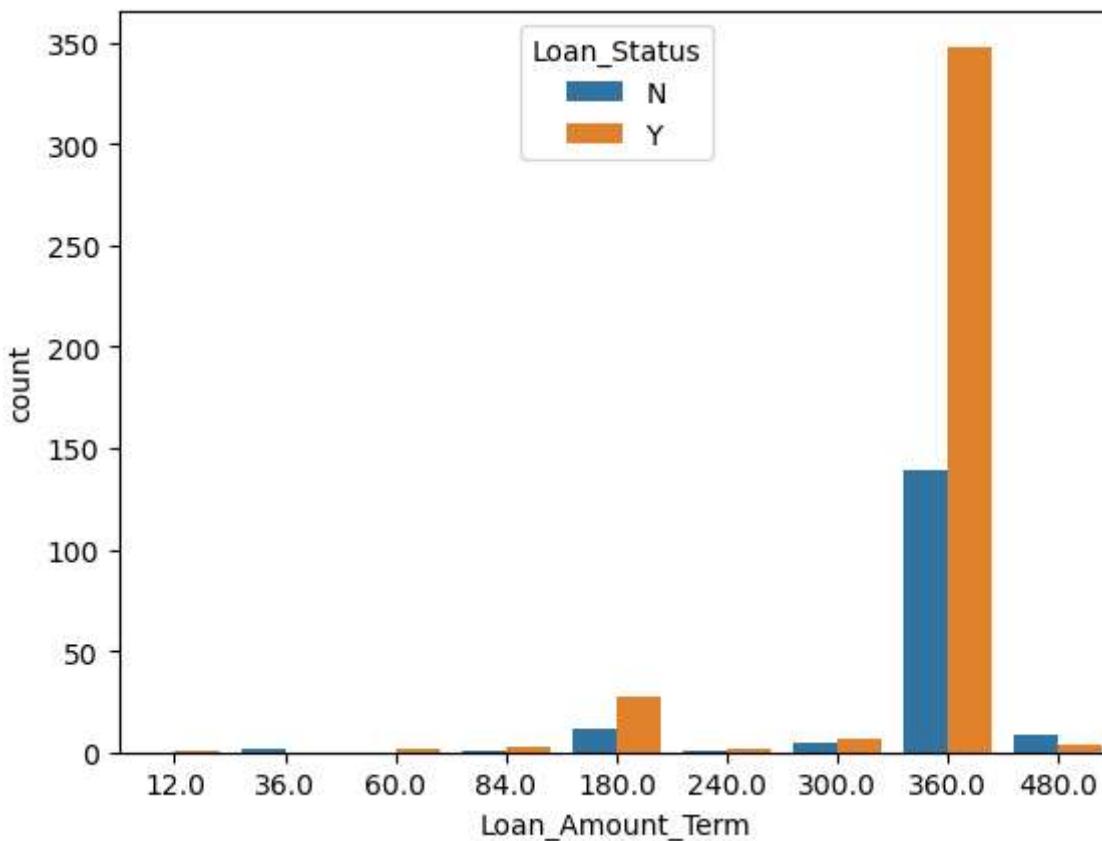
360.0	473
180.0	39
NaN	14
480.0	13
300.0	12
84.0	4
240.0	3
60.0	2
36.0	2
12.0	1

Name: Loan\_Amount\_Term, dtype: int64

```
In [44]: df['Loan_Amount_Term']=df['Loan_Amount_Term'].bfill()
```

```
In [45]: sns.countplot(x=df['Loan_Amount_Term'],hue=df['Loan_Status'])
```

```
Out[45]: <AxesSubplot:xlabel='Loan_Amount_Term', ylabel='count'>
```



```
In [46]: #columns in the dataset
df.columns
```

```
Out[46]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
       'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

```
In [47]: df['Credit_History'].value_counts(dropna=False, normalize=True)
```

```
Out[47]: 1.0    0.772647
0.0    0.143872
NaN    0.083481
Name: Credit_History, dtype: float64
```

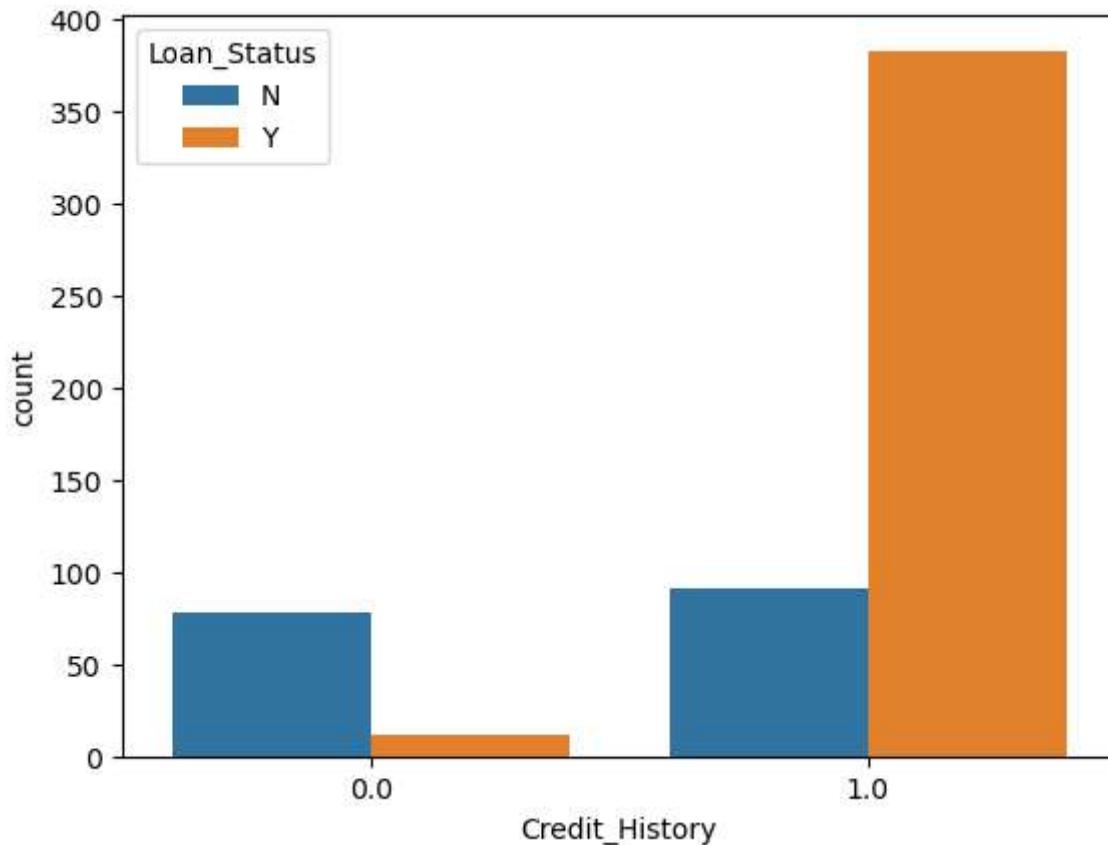
```
In [48]: df['Credit_History']=df['Credit_History'].bfill()
```

```
In [49]: df['Credit_History'].value_counts(dropna=False, normalize=True)
```

```
Out[49]: 1.0    0.841918
0.0    0.158082
Name: Credit_History, dtype: float64
```

```
In [50]: sns.countplot(x=df['Credit_History'], hue=df['Loan_Status'])
```

```
Out[50]: <AxesSubplot:xlabel='Credit_History', ylabel='count'>
```



creditscore having 1 have high chances of getting approval

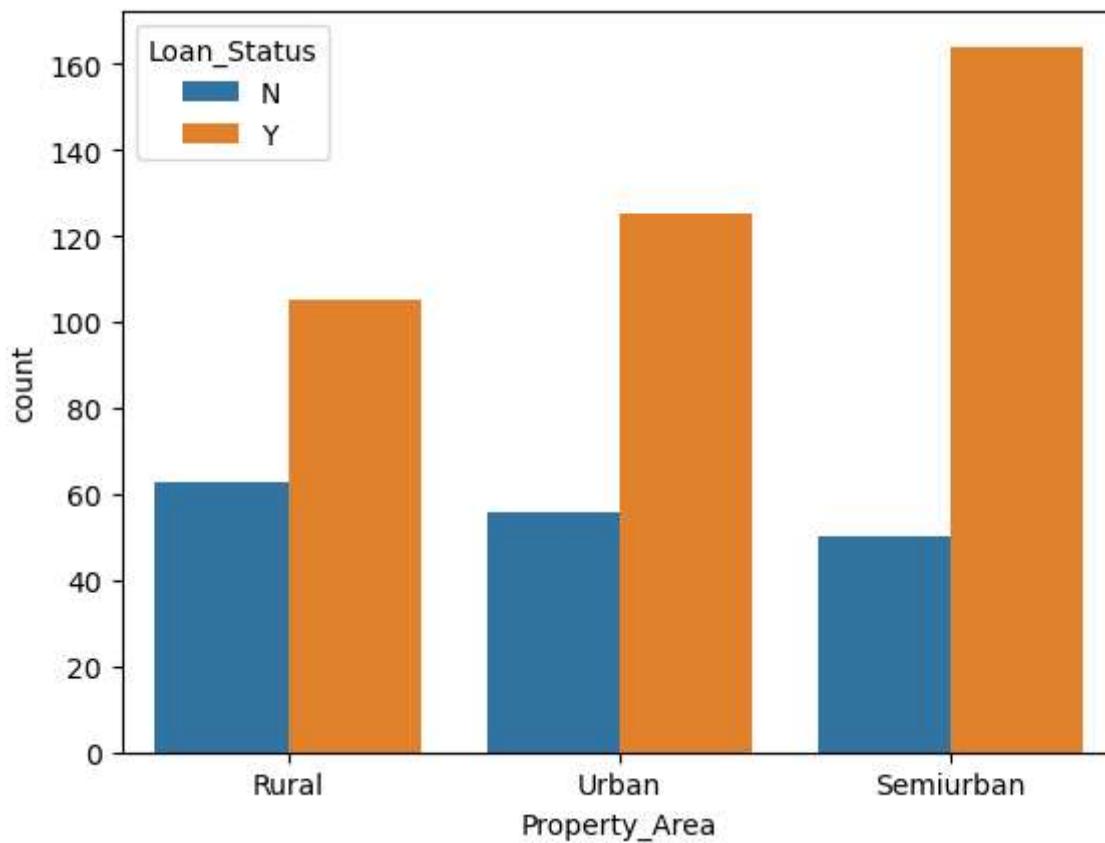
```
In [51]: df['Property_Area'].value_counts(dropna=False, normalize=True)
```

```
Out[51]:
```

Semiurban	0.380107
Urban	0.321492
Rural	0.298401
Name: Property_Area, dtype: float64	

```
In [52]: sns.countplot(x=df['Property_Area'], hue=df['Loan_Status'])
```

```
Out[52]: <AxesSubplot:xlabel='Property_Area', ylabel='count'>
```



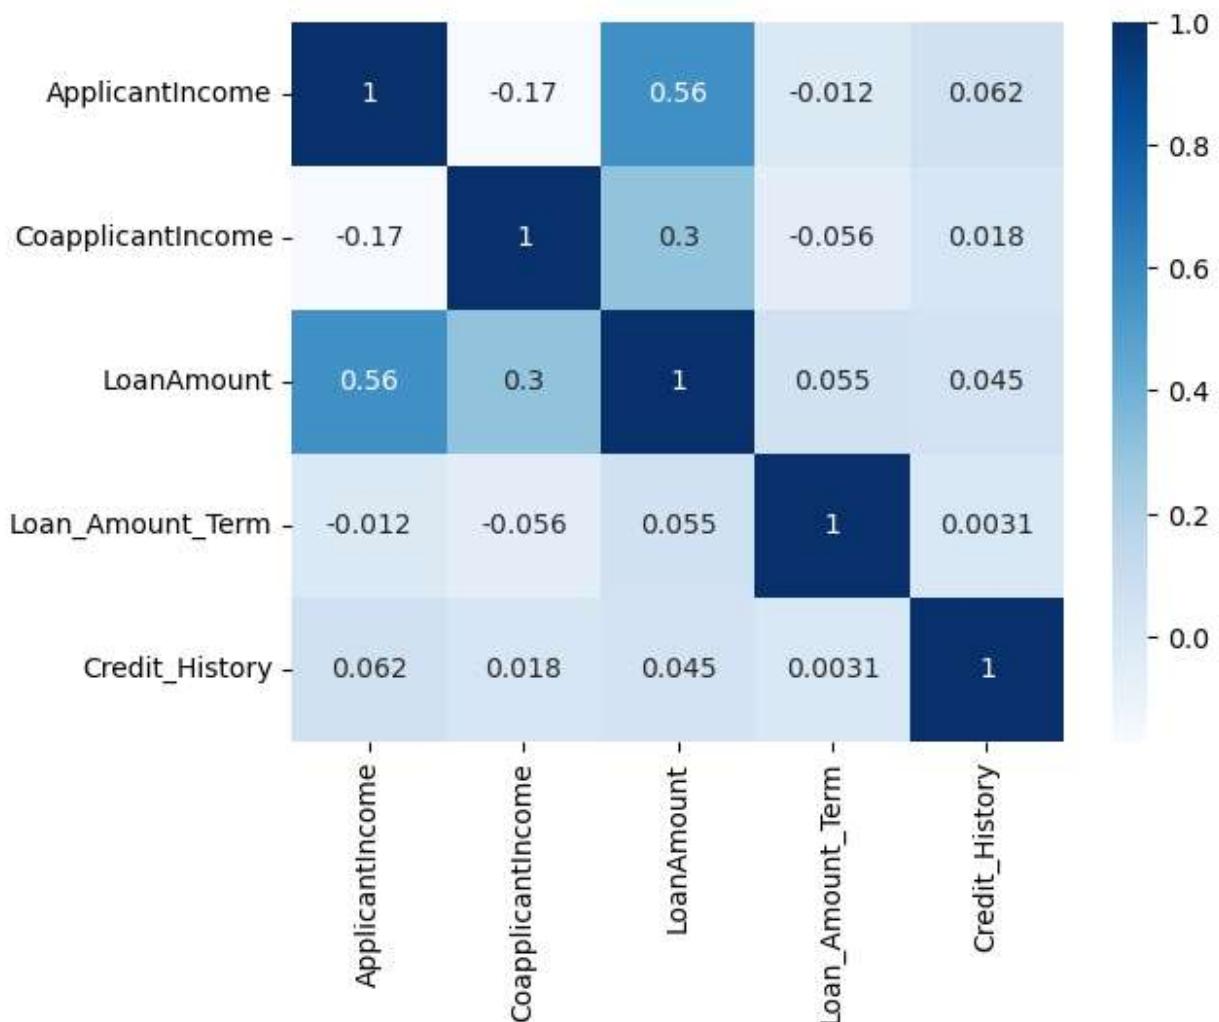
we can see people who are living in semiurban have high chance of getting loan approval

```
In [53]: df['LoanAmount'].value_counts(dropna=False, normalize=True)
```

```
Out[53]:
120.0    0.033748
110.0    0.028419
100.0    0.026643
160.0    0.021314
187.0    0.021314
...
89.0     0.001776
54.0     0.001776
78.0     0.001776
436.0    0.001776
253.0    0.001776
Name: LoanAmount, Length: 187, dtype: float64
```

```
In [54]: #Visualization of correlation using Heat map
sns.heatmap(df.corr(), annot=True, cmap='Blues')
```

```
Out[54]: <AxesSubplot:>
```



```
In [55]: #creating dummy
df.select_dtypes(include=object).columns
```

```
Out[55]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

```
In [56]: dummy=pd.get_dummies(df['Property_Area'],prefix='Property_Area')
dummy
```

Out[56]:

	Property_Area_Rural	Property_Area_Semiurban	Property_Area_Urban
<b>1</b>	1	0	0
<b>2</b>	0	0	1
<b>3</b>	0	0	1
<b>4</b>	0	0	1
<b>5</b>	0	0	1
...	...	...	...
<b>609</b>	1	0	0
<b>610</b>	1	0	0
<b>611</b>	0	0	1
<b>612</b>	0	0	1
<b>613</b>	0	1	0

563 rows × 3 columns

In [57]:

```
df=pd.concat([df,dummy],1)
df=df.drop('Property_Area',1)
df
```

C:\Users\gunnu\AppData\Local\Temp\ipykernel\_3032\2188141727.py:1: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only.

```
df=pd.concat([df,dummy],1)
```

C:\Users\gunnu\AppData\Local\Temp\ipykernel\_3032\2188141727.py:2: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only.

```
df=df.drop('Property_Area',1)
```

Out[57]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	Loan_Status
1	LP001003	Male	Yes	1	Graduate	No	4583		Approved
2	LP001005	Male	Yes	0	Graduate	Yes	3000		Approved
3	LP001006	Male	Yes	0	Not Graduate	No	2583		Approved
4	LP001008	Male	No	0	Graduate	No	6000		Approved
5	LP001011	Male	Yes	2	Graduate	Yes	5417		Approved
...	...	...	...	...	...	...	...	...	...
609	LP002978	Female	No	0	Graduate	No	2900		Approved
610	LP002979	Male	Yes	3	Graduate	No	4106		Approved
611	LP002983	Male	Yes	1	Graduate	No	8072		Approved
612	LP002984	Male	Yes	2	Graduate	No	7583		Approved
613	LP002990	Female	No	0	Graduate	Yes	4583		Approved

563 rows × 15 columns

In [58]:

```
dummy1=pd.get_dummies(df['Gender'],prefix='gender')
df=pd.concat([df,dummy1],1)
dummy2=pd.get_dummies(df['Education'],prefix='Education')
df=pd.concat([df,dummy2],1)
df
```

C:\Users\gunnu\AppData\Local\Temp\ipykernel\_3032\1941449558.py:2: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only.

```
df=pd.concat([df,dummy1],1)
```

C:\Users\gunnu\AppData\Local\Temp\ipykernel\_3032\1941449558.py:4: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only.

```
df=pd.concat([df,dummy2],1)
```

Out[58]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	Loan_Status
1	LP001003	Male	Yes	1	Graduate	No	4583		Approved
2	LP001005	Male	Yes	0	Graduate	Yes	3000		Approved
3	LP001006	Male	Yes	0	Not Graduate	No	2583		Approved
4	LP001008	Male	No	0	Graduate	No	6000		Approved
5	LP001011	Male	Yes	2	Graduate	Yes	5417		Approved
...	...	...	...	...	...	...	...	...	...
609	LP002978	Female	No	0	Graduate	No	2900		Approved
610	LP002979	Male	Yes	3	Graduate	No	4106		Approved
611	LP002983	Male	Yes	1	Graduate	No	8072		Approved
612	LP002984	Male	Yes	2	Graduate	No	7583		Approved
613	LP002990	Female	No	0	Graduate	Yes	4583		Approved

563 rows x 19 columns

In [59]: `#dropping the original columns after dummy creation  
d_col=['Gender','Education','Loan_ID']  
df.drop(d_col,axis=1,inplace=True)`

In [60]: `v = ['Loan_Status']  
def binary_map(x):  
 return x.map({'Y': 1, "N": 0})  
df[v] = df[v].apply(binary_map)`

In [61]: `varlist = ['Married','Self_Employed']  
def binary_map(x):  
 return x.map({'Yes': 1, "No": 0})  
df[varlist] = df[varlist].apply(binary_map)`

In [62]: `df`

Out[62]:

	Married	Dependents	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term
1	1	1	0	4583	1508.0	128.0	360
2	1	0	1	3000	0.0	66.0	360
3	1	0	0	2583	2358.0	120.0	360
4	0	0	0	6000	0.0	141.0	360
5	1	2	1	5417	4196.0	267.0	360
...	...	...	...	...	...	...	...
609	0	0	0	2900	0.0	71.0	360
610	1	3	0	4106	0.0	40.0	360
611	1	1	0	8072	240.0	253.0	360
612	1	2	0	7583	0.0	187.0	360
613	0	0	1	4583	0.0	133.0	360

563 rows × 16 columns

In [63]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 563 entries, 1 to 613
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Married          563 non-null    int64  
 1   Dependents       563 non-null    object 
 2   Self_Employed    563 non-null    int64  
 3   ApplicantIncome  563 non-null    int64  
 4   CoapplicantIncome 563 non-null    float64 
 5   LoanAmount        563 non-null    float64 
 6   Loan_Amount_Term  563 non-null    float64 
 7   Credit_History    563 non-null    float64 
 8   Loan_Status       563 non-null    int64  
 9   Property_Area_Rural 563 non-null    uint8  
 10  Property_Area_Semiurban 563 non-null    uint8  
 11  Property_Area_Urban  563 non-null    uint8  
 12  gender_Female     563 non-null    uint8  
 13  gender_Male       563 non-null    uint8  
 14  Education_Graduate 563 non-null    uint8  
 15  Education_Not Graduate 563 non-null    uint8  
dtypes: float64(4), int64(4), object(1), uint8(7)
memory usage: 47.8+ KB
```

In [64]: `df['Dependents']=pd.to_numeric(df['Dependents'])`  
`df['Credit_History']=pd.to_numeric(df['Credit_History'])`In [65]: `from sklearn.model_selection import train_test_split`  
`y = df['Loan_Status']`  
`X = df.drop('Loan_Status',axis=1)`

```
In [66]: X_train,X_test,y_train,y_test=train_test_split(X,y,train_size=0.7,random_state=50)
```

```
In [67]: X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 394 entries, 34 to 521
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Married          394 non-null    int64  
 1   Dependents       394 non-null    int64  
 2   Self_Employed    394 non-null    int64  
 3   ApplicantIncome  394 non-null    int64  
 4   CoapplicantIncome 394 non-null    float64 
 5   LoanAmount        394 non-null    float64 
 6   Loan_Amount_Term 394 non-null    float64 
 7   Credit_History   394 non-null    float64 
 8   Property_Area_Rural 394 non-null  uint8  
 9   Property_Area_Semiurban 394 non-null  uint8  
 10  Property_Area_Urban 394 non-null  uint8  
 11  gender_Female    394 non-null    uint8  
 12  gender_Male      394 non-null    uint8  
 13  Education_Graduate 394 non-null  uint8  
 14  Education_Not Graduate 394 non-null  uint8  
dtypes: float64(4), int64(4), uint8(7)
memory usage: 30.4 KB
```

```
In [68]: from sklearn.preprocessing import StandardScaler
Scaler = StandardScaler()
```

```
In [69]: num_cols =X_train.select_dtypes(include=['float64', 'int64']).columns
num_cols
```

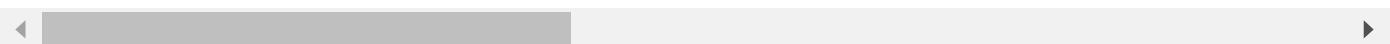
```
Out[69]: Index(['Married', 'Dependents', 'Self_Employed', 'ApplicantIncome',
 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term',
 'Credit_History'],
 dtype='object')
```

```
In [70]: X_train[num_cols] = Scaler.fit_transform(X_train[num_cols])
X_train
```

Out[70]:

	Married	Dependents	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan
34	-1.416910	2.208703	-0.402793	2.432051	0.918465	2.808206	
390	-1.416910	2.208703	-0.402793	1.369937	-0.863334	0.695737	
457	0.705761	-0.761362	-0.402793	-0.369661	0.662480	0.507962	
295	0.705761	2.208703	-0.402793	-0.144045	-0.120918	-0.477857	
482	0.705761	-0.761362	-0.402793	-0.887493	1.007555	-0.196195	
...	...	...	...	...	...	...	...
75	-1.416910	-0.761362	-0.402793	-0.356277	-0.863334	-0.430914	
148	-1.416910	-0.761362	-0.402793	1.635386	0.126158	1.321653	
315	0.705761	0.228660	-0.402793	-0.468128	0.110716	-0.462209	
121	-1.416910	-0.761362	-0.402793	-0.223712	-0.863334	-1.510620	
521	-1.416910	-0.761362	-0.402793	-0.754609	-0.863334	-1.338493	

394 rows × 15 columns

In [71]: `x_train.info()`

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 394 entries, 34 to 521
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Married          394 non-null    float64
 1   Dependents       394 non-null    float64
 2   Self_Employed    394 non-null    float64
 3   ApplicantIncome  394 non-null    float64
 4   CoapplicantIncome 394 non-null    float64
 5   LoanAmount        394 non-null    float64
 6   Loan_Amount_Term  394 non-null    float64
 7   Credit_History    394 non-null    float64
 8   Property_Area_Rural 394 non-null    uint8  
 9   Property_Area_Semiurban 394 non-null    uint8  
 10  Property_Area_Urban 394 non-null    uint8  
 11  gender_Female     394 non-null    uint8  
 12  gender_Male       394 non-null    uint8  
 13  Education_Graduate 394 non-null    uint8  
 14  Education_Not Graduate 394 non-null    uint8  
dtypes: float64(8), uint8(7)
memory usage: 30.4 KB

```

In [72]: `x_train.isnull().sum()`

```
Out[72]: Married          0
         Dependents      0
         Self_Employed    0
         ApplicantIncome   0
         CoapplicantIncome 0
         LoanAmount        0
         Loan_Amount_Term  0
         Credit_History     0
         Property_Area_Rural 0
         Property_Area_Semiurban 0
         Property_Area_Urban 0
         gender_Female      0
         gender_Male        0
         Education_Graduate 0
         Education_Not Graduate 0
         dtype: int64
```

```
In [73]: import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression()
```

```
In [74]: from sklearn.feature_selection import RFE
rfe = RFE(estimator=logreg, n_features_to_select=15) # running RFE with 15 variables
rfe = rfe.fit(X_train,y_train)
```

```
In [75]: col = X_train.columns[rfe.support_]
col
```

```
Out[75]: Index(['Married', 'Dependents', 'Self_Employed', 'ApplicantIncome',
       'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Credit_History',
       'Property_Area_Rural', 'Property_Area_Semiurban', 'Property_Area_Urban',
       'gender_Female', 'gender_Male', 'Education_Graduate',
       'Education_Not Graduate'],
      dtype='object')
```

```
In [76]: #BUILDING MODEL #1

X_train_sm = sm.add_constant(X_train[col])
logm1 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
res = logm1.fit()
res.summary()
```

Out[76]:

## Generalized Linear Model Regression Results

<b>Dep. Variable:</b>	Loan_Status	<b>No. Observations:</b>	394				
<b>Model:</b>	GLM	<b>Df Residuals:</b>	381				
<b>Model Family:</b>	Binomial	<b>Df Model:</b>	12				
<b>Link Function:</b>	Logit	<b>Scale:</b>	1.0000				
<b>Method:</b>	IRLS	<b>Log-Likelihood:</b>	-180.60				
<b>Date:</b>	Sun, 13 Aug 2023	<b>Deviance:</b>	361.20				
<b>Time:</b>	13:13:40	<b>Pearson chi2:</b>	396.				
<b>No. Iterations:</b>	12	<b>Pseudo R-squ. (CS):</b>	0.2622				
<b>Covariance Type:</b>	nonrobust						
		coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	0.3996	0.088	4.545	0.000	0.227	0.572	
<b>Married</b>	0.2617	0.151	1.728	0.084	-0.035	0.558	
<b>Dependents</b>	0.0572	0.147	0.388	0.698	-0.232	0.346	
<b>Self_Employed</b>	-0.0506	0.141	-0.360	0.719	-0.326	0.225	
<b>ApplicantIncome</b>	0.1236	0.192	0.645	0.519	-0.252	0.499	
<b>CoapplicantIncome</b>	0.0673	0.165	0.409	0.682	-0.255	0.390	
<b>LoanAmount</b>	-0.1932	0.187	-1.034	0.301	-0.559	0.173	
<b>Loan_Amount_Term</b>	0.0202	0.137	0.147	0.883	-0.249	0.290	
<b>Credit_History</b>	1.2422	0.147	8.440	0.000	0.954	1.531	
<b>Property_Area_Rural</b>	-0.1250	0.189	-0.663	0.507	-0.495	0.245	
<b>Property_Area_Semiurban</b>	0.5162	0.193	2.672	0.008	0.138	0.895	
<b>Property_Area_Urban</b>	0.0084	0.199	0.042	0.966	-0.381	0.398	
<b>gender_Female</b>	0.2339	0.222	1.054	0.292	-0.201	0.669	
<b>gender_Male</b>	0.1657	0.166	0.995	0.320	-0.161	0.492	
<b>Education_Graduate</b>	0.2295	0.148	1.553	0.121	-0.060	0.519	
<b>Education_Not Graduate</b>	0.1701	0.186	0.916	0.359	-0.194	0.534	

In [77]:

```
# Check for the VIF values of the feature variables.
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

In [78]:

```
# Create a dataframe that will contain the names of all the feature variables and their
vif = pd.DataFrame()
vif['Features'] = X_train[col].columns
vif['VIF'] = [variance_inflation_factor(X_train[col].values, i) for i in range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
C:\Users\gunnu\anaconda3\lib\site-packages\statsmodels\stats\outliers_influence.py:19
5: RuntimeWarning: divide by zero encountered in double_scalars
    vif = 1. / (1. - r_squared_i)
```

Out[78]:

	Features	VIF
8	Property_Area_Rural	inf
9	Property_Area_Semiurban	inf
10	Property_Area_Urban	inf
11	gender_Female	inf
12	gender_Male	inf
13	Education_Graduate	inf
14	Education_Not Graduate	inf
5	LoanAmount	1.93
3	ApplicantIncome	1.91
4	CoapplicantIncome	1.49
0	Married	1.34
1	Dependents	1.20
2	Self_Employed	1.12
6	Loan_Amount_Term	1.05
7	Credit_History	1.02

In [79]: `col=col.drop('Property_Area_Rural',1)`

```
#BUILDING MODEL #2
X_train_sm = sm.add_constant(X_train[col])
logm1 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
res = logm1.fit()
res.summary()
```

Out[80]:

## Generalized Linear Model Regression Results

<b>Dep. Variable:</b>	Loan_Status	<b>No. Observations:</b>	394				
<b>Model:</b>	GLM	<b>Df Residuals:</b>	381				
<b>Model Family:</b>	Binomial	<b>Df Model:</b>	12				
<b>Link Function:</b>	Logit	<b>Scale:</b>	1.0000				
<b>Method:</b>	IRLS	<b>Log-Likelihood:</b>	-180.60				
<b>Date:</b>	Sun, 13 Aug 2023	<b>Deviance:</b>	361.20				
<b>Time:</b>	13:13:40	<b>Pearson chi2:</b>	396.				
<b>No. Iterations:</b>	80	<b>Pseudo R-squ. (CS):</b>	0.2622				
<b>Covariance Type:</b>	nonrobust						
		coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	0.3371	0.135	2.503	0.012	0.073	0.601	
<b>Married</b>	0.2617	0.151	1.728	0.084	-0.035	0.558	
<b>Dependents</b>	0.0572	0.147	0.388	0.698	-0.232	0.346	
<b>Self_Employed</b>	-0.0506	0.141	-0.360	0.719	-0.326	0.225	
<b>ApplicantIncome</b>	0.1236	0.192	0.645	0.519	-0.252	0.499	
<b>CoapplicantIncome</b>	0.0673	0.165	0.409	0.682	-0.255	0.390	
<b>LoanAmount</b>	-0.1932	0.187	-1.034	0.301	-0.559	0.173	
<b>Loan_Amount_Term</b>	0.0202	0.137	0.147	0.883	-0.249	0.290	
<b>Credit_History</b>	1.2422	0.147	8.440	0.000	0.954	1.531	
<b>Property_Area_Semiurban</b>	0.6412	0.327	1.960	0.050	0.000	1.282	
<b>Property_Area_Urban</b>	0.1334	0.330	0.404	0.686	-0.514	0.781	
<b>gender_Female</b>	0.2027	0.232	0.873	0.383	-0.252	0.658	
<b>gender_Male</b>	0.1344	0.168	0.800	0.424	-0.195	0.464	
<b>Education_Graduate</b>	0.1982	0.159	1.248	0.212	-0.113	0.510	
<b>Education_Not Graduate</b>	0.1389	0.190	0.729	0.466	-0.234	0.512	

In [81]:

```
vif = pd.DataFrame()
vif['Features'] = X_train[col].columns
vif['VIF'] = [variance_inflation_factor(X_train[col].values, i) for i in range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

C:\Users\gunnu\anaconda3\lib\site-packages\statsmodels\stats\outliers\_influence.py:19  
 5: RuntimeWarning: divide by zero encountered in double\_scalars  
 vif = 1. / (1. - r\_squared\_i)

Out[81]:

	Features	VIF
10	gender_Female	inf
11	gender_Male	inf
12	Education_Graduate	inf
13	Education_Not Graduate	inf
5	LoanAmount	1.93
3	ApplicantIncome	1.91
4	CoapplicantIncome	1.49
8	Property_Area_Semiurban	1.41
9	Property_Area_Urban	1.41
0	Married	1.34
1	Dependents	1.20
2	Self_Employed	1.12
6	Loan_Amount_Term	1.05
7	Credit_History	1.02

In [82]:

```
#BUILDING MODEL #3
col=col.drop('gender_Female',1)

X_train_sm = sm.add_constant(X_train[col])
logm1 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
res = logm1.fit()
res.summary()
```

Out[82]:

## Generalized Linear Model Regression Results

<b>Dep. Variable:</b>	Loan_Status	<b>No. Observations:</b>	394			
<b>Model:</b>	GLM	<b>Df Residuals:</b>	381			
<b>Model Family:</b>	Binomial	<b>Df Model:</b>	12			
<b>Link Function:</b>	Logit	<b>Scale:</b>	1.0000			
<b>Method:</b>	IRLS	<b>Log-Likelihood:</b>	-180.60			
<b>Date:</b>	Sun, 13 Aug 2023	<b>Deviance:</b>	361.20			
<b>Time:</b>	13:13:40	<b>Pearson chi2:</b>	396.			
<b>No. Iterations:</b>	6	<b>Pseudo R-squ. (CS):</b>	0.2622			
<b>Covariance Type:</b>	nonrobust					
	coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	0.4722	0.267	1.769	0.077	-0.051	0.995
<b>Married</b>	0.2617	0.151	1.728	0.084	-0.035	0.558
<b>Dependents</b>	0.0572	0.147	0.388	0.698	-0.232	0.346
<b>Self_Employed</b>	-0.0506	0.141	-0.360	0.719	-0.326	0.225
<b>ApplicantIncome</b>	0.1236	0.192	0.645	0.519	-0.252	0.499
<b>CoapplicantIncome</b>	0.0673	0.165	0.409	0.682	-0.255	0.390
<b>LoanAmount</b>	-0.1932	0.187	-1.034	0.301	-0.559	0.173
<b>Loan_Amount_Term</b>	0.0202	0.137	0.147	0.883	-0.249	0.290
<b>Credit_History</b>	1.2422	0.147	8.440	0.000	0.954	1.531
<b>Property_Area_Semiurban</b>	0.6412	0.327	1.960	0.050	0.000	1.282
<b>Property_Area_Urban</b>	0.1334	0.330	0.404	0.686	-0.514	0.781
<b>gender_Male</b>	-0.0682	0.382	-0.179	0.858	-0.817	0.681
<b>Education_Graduate</b>	0.2658	0.188	1.415	0.157	-0.102	0.634
<b>Education_Not Graduate</b>	0.2064	0.230	0.898	0.369	-0.244	0.657

In [83]:

```
vif = pd.DataFrame()
vif['Features'] = X_train[col].columns
vif['VIF'] = [variance_inflation_factor(X_train[col].values, i) for i in range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[83]:

	Features	VIF
11	Education_Graduate	7.36
12	Education_Not Graduate	3.08
5	LoanAmount	1.93
3	ApplicantIncome	1.91
4	CoapplicantIncome	1.49
8	Property_Area_Semiurban	1.41
9	Property_Area_Urban	1.41
0	Married	1.34
10	gender_Male	1.24
1	Dependents	1.20
2	Self_Employed	1.12
6	Loan_Amount_Term	1.05
7	Credit_History	1.02

In [84]:

```
#BUILDING MODEL #3
col=col.drop('Education_Graduate',1)

X_train_sm = sm.add_constant(X_train[col])
logm1 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
res = logm1.fit()
res.summary()
```

Out[84]:

## Generalized Linear Model Regression Results

<b>Dep. Variable:</b>	Loan_Status	<b>No. Observations:</b>	394				
<b>Model:</b>	GLM	<b>Df Residuals:</b>	381				
<b>Model Family:</b>	Binomial	<b>Df Model:</b>	12				
<b>Link Function:</b>	Logit	<b>Scale:</b>	1.0000				
<b>Method:</b>	IRLS	<b>Log-Likelihood:</b>	-180.60				
<b>Date:</b>	Sun, 13 Aug 2023	<b>Deviance:</b>	361.20				
<b>Time:</b>	13:13:40	<b>Pearson chi2:</b>	396.				
<b>No. Iterations:</b>	5	<b>Pseudo R-squ. (CS):</b>	0.2622				
<b>Covariance Type:</b>	nonrobust						
		coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	0.7380	0.400	1.844	0.065	-0.047	1.523	
<b>Married</b>	0.2617	0.151	1.728	0.084	-0.035	0.558	
<b>Dependents</b>	0.0572	0.147	0.388	0.698	-0.232	0.346	
<b>Self_Employed</b>	-0.0506	0.141	-0.360	0.719	-0.326	0.225	
<b>ApplicantIncome</b>	0.1236	0.192	0.645	0.519	-0.252	0.499	
<b>CoapplicantIncome</b>	0.0673	0.165	0.409	0.682	-0.255	0.390	
<b>LoanAmount</b>	-0.1932	0.187	-1.034	0.301	-0.559	0.173	
<b>Loan_Amount_Term</b>	0.0202	0.137	0.147	0.883	-0.249	0.290	
<b>Credit_History</b>	1.2422	0.147	8.440	0.000	0.954	1.531	
<b>Property_Area_Semiurban</b>	0.6412	0.327	1.960	0.050	0.000	1.282	
<b>Property_Area_Urban</b>	0.1334	0.330	0.404	0.686	-0.514	0.781	
<b>gender_Male</b>	-0.0682	0.382	-0.179	0.858	-0.817	0.681	
<b>Education_Not Graduate</b>	-0.0593	0.324	-0.183	0.855	-0.694	0.575	

In [85]:

```
vif = pd.DataFrame()
vif['Features'] = X_train[col].columns
vif['VIF'] = [variance_inflation_factor(X_train[col].values, i) for i in range(X_train.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

Out[85]:

	Features	VIF
10	gender_Male	2.75
5	LoanAmount	1.93
3	ApplicantIncome	1.91
8	Property_Area_Semiurban	1.70
9	Property_Area_Urban	1.69
4	CoapplicantIncome	1.48
11	Education_Not Graduate	1.38
0	Married	1.25
1	Dependents	1.19
2	Self_Employed	1.12
6	Loan_Amount_Term	1.05
7	Credit_History	1.02

In [86]: 

```
y_train_pred=res.predict(X_train_sm)
y_train_pred[:10]
```

Out[86]:

```
34    0.695057
390   0.727285
457   0.805940
295   0.853134
482   0.883600
552   0.823216
329   0.751549
3     0.813366
497   0.824725
389   0.802651
dtype: float64
```

In [87]: 

```
y_train_pred=y_train_pred.values.reshape(-1)
y_train_pred[:10]
```

Out[87]:

```
array([0.69505722, 0.7272854 , 0.80593984, 0.853134 , 0.88359987,
       0.82321638, 0.75154945, 0.81336584, 0.82472469, 0.80265075])
```

In [88]: 

```
y_train_pred[:10]
```

Out[88]:

```
array([0.69505722, 0.7272854 , 0.80593984, 0.853134 , 0.88359987,
       0.82321638, 0.75154945, 0.81336584, 0.82472469, 0.80265075])
```

In [89]: 

```
y_train_pred_final=pd.DataFrame({'Status':y_train.values,'Status_prob':y_train_pred})
y_train_pred_final['Prospect ID'] = y_train.index
y_train_pred_final
```

Out[89]:

	Status	Status_prob	Prospect ID
<b>0</b>	0	0.695057	34
<b>1</b>	1	0.727285	390
<b>2</b>	0	0.805940	457
<b>3</b>	1	0.853134	295
<b>4</b>	1	0.883600	482
...	...	...	...
<b>389</b>	0	0.728456	75
<b>390</b>	0	0.702170	148
<b>391</b>	1	0.817917	315
<b>392</b>	1	0.851934	121
<b>393</b>	1	0.829596	521

394 rows × 3 columns

```
In [90]: y_train_pred_final['Predicted'] = y_train_pred_final.Status_prob.map(lambda x: 1 if x  
# Let's see the head  
y_train_pred_final.head(25)
```

Out[90]:

	Status	Status_prob	Prospect ID	Predicted
0	0	0.695057	34	1
1	1	0.727285	390	1
2	0	0.805940	457	1
3	1	0.853134	295	1
4	1	0.883600	482	1
5	1	0.823216	552	1
6	1	0.751549	329	1
7	1	0.813366	3	1
8	1	0.824725	497	1
9	1	0.802651	389	1
10	1	0.838602	394	1
11	1	0.131685	411	0
12	1	0.802965	441	1
13	0	0.094444	412	0
14	1	0.889715	560	1
15	1	0.836452	542	1
16	0	0.097497	396	0
17	1	0.812971	68	1
18	1	0.820237	176	1
19	1	0.904182	335	1
20	1	0.722824	197	1
21	1	0.828069	249	1
22	0	0.226352	7	0
23	0	0.147294	569	0
24	1	0.884396	234	1

In [91]:

```
from sklearn import metrics

# Confusion matrix
confusion = metrics.confusion_matrix(y_train_pred_final.Status, y_train_pred_final.Predicted)
print(confusion)

[[ 56  62]
 [  9 267]]
```

In [92]:

```
# Let's check the overall accuracy.
```

```
print(metrics.accuracy_score(y_train_pred_final.Status, y_train_pred_final.Predicted))
```

0.8197969543147208

```
In [93]: TP = confusion[1,1] # true positive
          TN = confusion[0,0] # true negatives
          FP = confusion[0,1] # false positives
          FN = confusion[1,0] # false negatives
```

```
In [94]: # Lets see the sensitivity of our logistic regression model
          TP / float(TP+FN)
```

```
Out[94]: 0.967391304347826
```

```
In [95]: # Let us calculate specificity
          TN / float(TN+FP)
```

```
Out[95]: 0.4745762711864407
```

```
In [96]: # Calculate False Positive Rate - predicting conversion when customer does not have cor
          print(FP / float(TN+FP))
```

```
0.5254237288135594
```

```
In [97]: print (TP / float(TP+FP))#positive predictive value
```

```
0.8115501519756839
```

```
In [98]: # Negative predictive value
          print (TN / float(TN+ FN))
```

```
0.8615384615384616
```

```
In [99]: # prediction test
```

```
In [100...]: num_cols=X_test.select_dtypes(include=['float64','int64']).columns
          num_cols
```

```
Out[100]: Index(['Married', 'Dependents', 'Self_Employed', 'ApplicantIncome',
       'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term',
       'Credit_History'],
      dtype='object')
```

```
In [101...]: X_test[num_cols]=Scaler.transform(X_test[num_cols])
          X_test=X_test[col]
          X_test
```

Out[101]:

	Married	Dependents	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan
240	0.705761	1.218681	-0.402793	0.303043	2.106331	-0.321378	
123	0.705761	1.218681	-0.402793	-0.608979	-0.863334	-0.931647	
310	-1.416910	-0.761362	-0.402793	-0.621725	-0.863334	-0.884703	
227	0.705761	1.218681	-0.402793	0.440388	0.143382	1.086935	
502	0.705761	1.218681	-0.402793	-0.000964	2.476946	1.055639	
...	...	...	...	...	...	...	...
23	0.705761	1.218681	-0.402793	-0.478963	0.275236	-0.446561	
110	-1.416910	-0.761362	-0.402793	0.142754	-0.863334	-0.071011	
530	0.705761	-0.761362	-0.402793	-1.224641	2.403298	1.180822	
406	0.705761	-0.761362	-0.402793	-0.860725	0.562105	-0.399618	
511	0.705761	0.228660	-0.402793	0.381435	0.326908	1.712851	

169 rows × 12 columns

In [102...]

```
X_test_sm=sm.add_constant(X_test)
y_test_pred=res.predict(X_test_sm)
y_test_pred.head()
```

Out[102]:

```
240    0.851226
123    0.899425
310    0.829168
227    0.880611
502    0.891563
dtype: float64
```

In [103...]

```
y_test_pred=y_test_pred.values.reshape(-1)
```

In [104...]

```
# Converting y_pred to a dataframe which is an array
y_pred_final = pd.DataFrame({'Status':y_test.values,'Status_prob':y_test_pred})
y_pred_final['Prospect ID'] = y_test.index
```

In [105...]

```
# Let's see the head of y_pred_final
y_pred_final.head()
```

Out[105]:

	Status	Status_prob	Prospect ID
0	1	0.851226	240
1	1	0.899425	123
2	1	0.829168	310
3	1	0.880611	227
4	1	0.891563	502

```
In [106]: y_pred_final['final_predicted'] = y_pred_final.Status_prob.map(lambda x: 1 if x > 0.3  
y_pred_final.head()
```

Out[106]:

	Status	Status_prob	Prospect ID	final_predicted
<b>0</b>	1	0.851226	240	1
<b>1</b>	1	0.899425	123	1
<b>2</b>	1	0.829168	310	1
<b>3</b>	1	0.880611	227	1
<b>4</b>	1	0.891563	502	1

```
In [107]: metrics.accuracy_score(y_pred_final.Status, y_pred_final.final_predicted)
```

Out[107]: 0.8165680473372781

```
In [108]: confusion2 = metrics.confusion_matrix(y_pred_final.Status, y_pred_final.final_predicted)  
confusion2
```

Out[108]: array([[ 22, 29],  
 [ 2, 116]], dtype=int64)

```
In [109]: TP = confusion2[1,1] # true positive  
TN = confusion2[0,0] # true negatives  
FP = confusion2[0,1] # false positives  
FN = confusion2[1,0] # false negatives
```

```
In [110]: #Let's see the sensitivity of our logistic regression model  
TP / float(TP+FN)
```

Out[110]: 0.9830508474576272

```
In [111]: TN / float(TN+FP)  
#for caculation specificity
```

Out[111]: 0.43137254901960786

```
In [112]: metrics.precision_score(y_pred_final.Status , y_pred_final.final_predicted)
```

Out[112]: 0.8

```
In [113]: metrics.recall_score(y_pred_final.Status, y_pred_final.final_predicted)
```

Out[113]: 0.9830508474576272

In [ ]: